Benchmarking Factual Accuracy of Large Language Models Across Scientific Domains

Abstract
The rapid proliferation and architectural scaling of Large Language Models (LLMs) have fundamentally transformed computational linguistics, natural language processing, and the paradigms of automated scientific discovery. Despite their exceptional capabilities in text generation, summarization, and multi-step reasoning, LLMs remain fundamentally constrained by the phenomenon of hallucination—the generation of plausible yet factually incorrect, logically inconsistent, or empirically unsupported content. In high-stakes scientific domains such as clinical medicine, synthetic chemistry, and theoretical physics, factual inaccuracy poses severe risks to research integrity, physical safety, and real-world application. This paper provides an exhaustive examination of the frameworks, methodologies, and datasets utilized to benchmark the factual accuracy of LLMs across diverse scientific disciplines. By systematically reviewing specialized benchmarks—including GPQA, SciEval, ChemBench, Med-PaLM, and TruthfulQA—the analysis explores the nuances of evaluating intrinsic and extrinsic factuality, the challenges of multimodal scientific reasoning, and the epistemological problem of "scalable oversight" when artificial intelligence capabilities surpass human expert verification. Furthermore, this research synthesizes current approaches to factuality evaluation, such as atomic fact verification and representation-space detection, highlights critical limitations including data contamination and inverse scaling, and delineates future trajectories for developing reliable, domain-specific artificial intelligence systems.
Keywords: Large Language Models (LLMs), Factual Accuracy, Hallucination, Scientific Benchmarking, Scalable Oversight, Computational Linguistics, Artificial Intelligence Evaluation, Retrieval-Augmented Generation (RAG).

Introduction

The integration of Large Language Models (LLMs) into the scientific workflow represents a paradigm shift in how empirical knowledge is synthesized, hypotheses are generated, and complex data is analyzed[1]. Models such as OpenAI's GPT-4, Anthropic's Claude 3.5, and highly specialized, domain-adapted iterations like Google's Med-PaLM 2 have demonstrated extraordinary proficiency in tasks ranging from clinical diagnostic reasoning to complex chemical synthesis prediction and graduate-level physics problem-solving[3]. However, the foundational architecture of these models—predicting the next token based on statistical distributions learned from vast, often uncurated corpora—renders them inherently susceptible to generating fabrications. This phenomenon, widely termed "hallucination," involves the production of content that is grammatically fluent and contextually plausible but fundamentally ungrounded in empirical reality, established scientific consensus, or the specific context provided by the user[6].

In the context of creative writing, marketing, or general conversational agents, factual deviations may be viewed as benign artifacts of generative flexibility or acceptable margins of error. Conversely, in specialized scientific and medical domains, the tolerance for factual inaccuracy is practically zero. A hallucinated drug interaction, a fabricated chemical property, an incorrect mathematical proof, or a misapplied physical theorem can lead to severe real-world consequences, undermining the utility of artificial intelligence in critical infrastructure and healthcare[6]. As models are increasingly deployed to summarize electronic health records, predict material stability, and assist in peer review, the assurance of absolute factual precision is imperative.

Consequently, the development of robust, rigorous, and domain-specific benchmarking frameworks has become a paramount priority within the artificial intelligence research community. This paper delivers a comprehensive thematic and technical analysis of the state-of-the-art in benchmarking LLM factual accuracy across scientific domains. It dissects the taxonomy of hallucinations, reviews pivotal literature and datasets, and evaluates the distinct challenges of measuring truthfulness in fields requiring advanced mathematical computation, deep domain expertise, and complex logical reasoning. Ultimately, the discussion elucidates how the science of benchmarking is evolving to address the complex reality of models that can increasingly outpace the verification capabilities of human non-experts, necessitating novel paradigms of machine evaluation and algorithmic oversight.

Background, Related Work, and Literature Review

To establish a foundation for understanding the evaluation of factual accuracy, it is essential to trace the scholarly literature that defines, categorizes, and measures hallucinations in generative models. Because the integration of exhaustive bibliographic information was uniquely requested for this analysis, Table 1 presents a synthesis of the foundational scholarly literature anchoring this discourse, embedding the requisite metadata directly into the analytical framework rather than appending an isolated reference section.

Table 1: Foundational Literature and Bibliographic Metadata in Factuality Benchmarking
Authors & Title
Rawte, V., Sheth, A., Das, A. *A Survey of Hallucination in Large Foundation Models*
Huang, L., et al. *A Survey on Hallucination in Large Language Models: Principles, Taxonomy,*

Challenges...
Singhal, K., et al. *Towards Expert-Level Medical Question Answering with Large Language Models*
Lin, S., Hilton, J., Evans, O. *TruthfulQA: Measuring How Models Mimic Human Falsehoods*
Min, S., et al. *FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation*
Rein, D., et al. *GPQA: A Graduate-Level Google-Proof Q&A Benchmark*
Jin, Q., et al. *PubMedQA: A Dataset for Biomedical Research Question Answering*
Sun, L., et al. *SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research*
Jablonka, K.M., et al. *Are large language models superhuman chemists? (ChemBench)*

The Taxonomy and Mechanics of Hallucination in LLMs

To rigorously benchmark factual accuracy, it is first necessary to delineate the exact nature of the inaccuracies being measured. Recent comprehensive surveys by Rawte et al. (2023) and Huang et al. (2023) have formalized the taxonomy of LLM hallucinations into distinct, measurable categories, moving the discipline beyond vague definitions of AI error[2].

Hallucinations are broadly classified along two primary axes: the origin of the information and the nature of the error. The first axis distinguishes between intrinsic and extrinsic hallucinations. Intrinsic hallucinations occur when the generated text explicitly contradicts the source material provided in the prompt, such as an LLM summarizing a clinical trial report but reversing the statistical outcomes[2]. Extrinsic hallucinations occur when the model introduces supplementary information that cannot be verified by the provided context. While extrinsic hallucinations are not necessarily false in the real world, they represent a failure of faithfulness to the input and introduce significant risk in closed-domain tasks[2].

The second axis differentiates factual from semantic hallucinations. Semantic hallucinations involve inconsistencies in logic, linguistic coherence, or direct contradictions within the model's own generated text, rendering the output nonsensical[2]. Factual hallucinations, which constitute the primary concern of scientific benchmarking, involve direct contradictions of established real-world knowledge[2]. Factual hallucinations are further divided by Huang et al. into *intrinsic factuality hallucinations*—errors rooted in flawed internal parametric knowledge representations—and *extrinsic factuality hallucinations*—failures to correctly retrieve, process, or integrate external knowledge when utilizing tools like Retrieval-Augmented Generation (RAG)[6].

The etiology of these errors is profoundly multifaceted. Data-related causes heavily influence hallucination rates; training on uncurated corpora containing factual errors, outdated information, or inherent biases directly embeds falsehoods into the model's weights[6]. Furthermore, source-reference divergence in training datasets—where reference summaries contain claims unsupported by their source documents—teaches the model to extrapolate beyond the evidence[2]. Model-specific factors, including the architectural constraints of attention mechanisms in processing long-context windows and the inherent probabilistic nature of decoding strategies (such as temperature scaling and top-p sampling), also significantly contribute to the generation of falsehoods[2]. When forced to predict the next token, the model prioritizes statistical likelihood and grammatical fluency over epistemological truth.

The Evolution of Factuality Evaluation Metrics

Historically, natural language generation was evaluated using superficial lexical overlap metrics such as BLEU and ROUGE. These metrics, while historically useful for rudimentary machine translation, are entirely inadequate for assessing factual consistency in science. A generated sentence can share immense lexical overlap with a reference text while negating its entire meaning through the insertion of a single modifier[24].

The literature reflects a decisive shift toward fact-based and classifier-based metrics. Fact-based detection extracts discrete claims from the LLM output and verifies them against established knowledge bases using exact matching or embedding similarity[6]. Classifier-based metrics employ specialized natural language inference (NLI) models to calculate an entailment score, determining whether the model's output is logically entailed by, contradicts, or is neutral toward a given scientific premise[6]. Furthermore, Question-Answering (QA) based metrics evaluate consistency by generating questions from the output and verifying if the model provides consistent answers across multiple iterations, while uncertainty estimation techniques analyze the model's internal confidence signals to flag potential fabrications[2].

The literature conclusively demonstrates that measuring factual accuracy in general domains (e.g., historical dates, geographical locations) is vastly insufficient for evaluating the utility of LLMs in scientific research. Scientific factuality requires the measurement of complex spatial reasoning, multi-step algorithmic calculation, adherence to physical laws, and the synthesis of contradictory literature. This necessity has driven the rapid creation of domain-specific benchmarks that test not merely the retrieval of static facts, but the rigorous application of those facts to novel, dynamic scenarios[6].

Main Thematic and Technical Discussion: Benchmarking Across Domains

The landscape of scientific benchmarking is highly segmented, with distinct methodologies, datasets, and philosophical approaches applied to general scientific reasoning, medicine, chemistry, and structured knowledge verification.

General Science, Mathematics, and Graduate-Level Reasoning

Evaluating an LLM's comprehension of fundamental scientific principles requires frameworks that test varying levels of cognitive complexity, distinguishing mere memorization from genuine analytical synthesis.

SciEval and the Application of Bloom's Taxonomy

Introduced by Sun et al. in 2023, **SciEval** is a comprehensive, multi-disciplinary evaluation benchmark encompassing approximately 18,000 rigorous questions across the fundamental scientific fields of physics, chemistry, and biology[22]. The benchmark's architectural foundation is derived from Bloom's Taxonomy, structuring the evaluation across four progressive cognitive dimensions to mirror human learning: basic knowledge, knowledge application, scientific calculation, and advanced research ability[25].

A critical, pervasive vulnerability of static benchmarks is the phenomenon of "data leakage"—the high probability that standardized test questions were inadvertently included in the LLM's massive training corpus[21]. When evaluating on leaked data, researchers are measuring data retrieval and memorization rather than inferential reasoning. To systematically counter this, SciEval introduced a paradigm-shifting "dynamic" data subset. These questions are algorithmically generated based on underlying scientific principles, ensuring the exact formulations have never appeared in pre-training data[25].

The empirical findings from SciEval were deeply revelatory regarding the current state of artificial intelligence. While top-tier models like GPT-4 achieved a highly impressive average accuracy of 73.93% on static data, their performance degraded precipitously on dynamic problems[25]. Specifically, models consistently struggled with dynamic scientific calculations in physics, often performing near random chance levels, revealing fundamental limitations in applying mathematical principles to novel problem-solving contexts[25]. Furthermore, while top models demonstrated competence in understanding experimental design, they failed significantly in analyzing experimental results and drawing meaningful conclusions, sharply delineating the boundary between a model's capacity for knowledge recall and its capacity for genuine scientific deduction[25].

GPQA: The Epistemological Challenge of Scalable Oversight

Perhaps the most conceptually ambitious benchmark in recent literature is **GPQA (Graduate-Level Google-Proof Q&A Benchmark)**, developed by Rein, Hilton, Evans, Michael, and colleagues[14]. GPQA consists of 448 exceptionally difficult multiple-choice questions in biology, physics, and chemistry, meticulously crafted by domain experts holding or actively pursuing PhDs in their respective fields[29].

The defining characteristic of GPQA is its explicitly "Google-proof" nature, designed to resist shallow web searches and basic information retrieval. During the rigorous validation phase, highly skilled non-expert validators—individuals possessing PhDs in unrelated disciplines—were given unrestricted internet access and spent an average of over 30 minutes attempting to solve each question[17]. Despite high intelligence, extensive education, and unlimited digital resources, these validators achieved only a 34% accuracy rate[17]. In stark contrast, true domain experts achieved 65% accuracy (rising to 74% when accounting for clear mistakes identified in retrospect), while the baseline GPT-4 model achieved 39%[29]. By March 2024, newer iterations such as Claude 3 Opus reportedly achieved approximately 60%, signaling rapid capability advancement[30].

GPQA is not merely a test of factual accuracy; it is an instrument designed to research *scalable oversight*[29]. Scalable oversight addresses a core AI safety dilemma: how can human operators reliably supervise, verify, and extract truthful information from AI systems when those systems possess capabilities that vastly surpass human verifiability? Because GPQA questions are impervious to superficial fact-checking by non-experts, they serve as a highly realistic testbed for evaluating advanced oversight mechanisms—such as AI safety via debate, market-making, and recursive reward modeling—which will be necessary to manage future AI systems that generate novel scientific knowledge[29].

SciBench: College-Level Problem Solving

Addressing higher education and complex mathematical paradigms, Wang et al. (2023) developed **SciBench**, which focuses explicitly on college-level scientific problem-solving[31]. SciBench comprises 695 advanced, open-ended problems sourced from rigorous instructional textbooks across mathematics and physics[31]. The benchmark specifically targets complex reasoning capabilities and advanced calculation skills, forcing models to utilize domain knowledge in conjunction with continuous mathematical formulation. Performance on SciBench consistently underscores the limitations of current LLMs in maintaining logical coherence and numerical precision through protracted, multi-step mathematical derivations, a common failure point that general knowledge benchmarks fail to expose[31].

The Medical and Biomedical Domain

The integration of LLMs into clinical healthcare and biomedical research necessitates the absolute highest stringency in factual accuracy, given the direct, immediate implications for patient safety, morbidity, and mortality.

Med-PaLM and the Evolution of Clinical QA

The benchmark standard in clinical knowledge evaluation has heavily revolved around the United States Medical Licensing Examination (USMLE), operationalized in datasets like MedQA[1]. The development of **Med-PaLM** (Singhal et al., *Nature*, 2023) marked a watershed moment in medical artificial intelligence, as it was the first model to exceed the USMLE passing threshold,

achieving a score of 67.2% on the MedQA dataset[1].
Recognizing that multiple-choice testing is insufficient for clinical readiness, Google Research subsequently developed **Med-PaLM 2**, which leveraged the highly advanced PaLM 2 base architecture, extensive medical domain fine-tuning, and sophisticated prompting strategies such as *ensemble refinement* (ER) and chain of retrieval[1]. Med-PaLM 2 achieved an unprecedented 86.5% on the MedQA dataset, improving upon its predecessor by over 19%[1]. It also demonstrated near state-of-the-art performance across MedMCQA (72.3%), PubMedQA (75.0%), and the MMLU clinical topics dataset[1].
Crucially, the evaluation of Med-PaLM 2 extended far beyond automated multiple-choice scoring into the realm of qualitative clinical utility. In a rigorous pairwise comparative ranking of long-form responses to 1,066 consumer medical questions, independent physician raters preferred Med-PaLM 2 answers over those generated by human physicians on eight out of nine critical axes, including factuality, medical reasoning capability, and the reflection of medical consensus (preferred 72.9% of the time)[1]. Furthermore, the researchers introduced datasets comprising 240 long-form "adversarial" questions specifically designed to probe LLM limitations, biases, and safety vulnerabilities[1]. Against these adversarial probes, Med-PaLM 2 was rated as having a low risk of harm for 90.6% of its answers, significantly outperforming earlier iterations[1].

PubMedQA and Biomedical Reasoning

For academic biomedical research, **PubMedQA** (Jin et al., 2019) remains a foundational benchmark, evaluating models on their ability to answer specific research questions based strictly on PubMed abstracts[19]. The dataset consists of 1,000 expert-annotated instances, 61,200 unlabeled instances, and 211,300 artificially generated QA instances[19]. By constraining the answers to yes, no, or maybe, the dataset forces models to deeply reason over quantitative contents, statistical significance, and complex biomedical terminology found in the context abstract, rather than generating open-ended text[19]. The performance gap observed during its introduction—where BioBERT achieved 68.1% accuracy compared to a single human expert's 78.0%—illustrates the profound difficulty of synthesizing abstractive conclusions from dense, empirical biomedical literature[19].

Multi-Criteria Decision Analysis in Healthcare Annotation

Evaluating open-ended clinical data extraction requires nuanced frameworks that move beyond binary correctness. Nemati et al. (2025) proposed a rigorous evaluation framework tailored for unstructured healthcare records, leveraging Multi-Criteria Decision Analysis (MCDA)[36]. The methodology structures performance scores across various quantitative and qualitative metrics into a unified decision matrix, standardizes them to ensure comparability, and identifies the Positive Ideal Solution (PIS) and Negative Ideal Solution (NIS) to rank models[36]. Metrics are weighted meticulously: human validation holds the highest weight, followed closely by Precision, Recall, and F1-Score, while semantic textual similarity, factual consistency, coherence, and fluency are assigned distinct weights to ensure balance[36]. This granular, mathematical approach to evaluation highlights that factuality in clinical records is a complex composite of precision, narrative consistency, and clinical relevance, rather than a simple metric of lexical overlap.

Hallucination Testing via Med-HALT

To directly combat the risks of fabrication in healthcare, specialized diagnostic benchmarks like **Med-HALT** (Medical Domain Hallucination Test, Umapathi et al., 2023) have been developed explicitly to trigger, evaluate, and subsequently reduce hallucinations in medical LLMs[37]. By testing leading models (e.g., GPT-3.5, Llama-2, MPT, and Falcon) under adversarial medical conditions, Med-HALT reveals stark performance disparities and provides a standardized methodology for auditing model safety before clinical deployment[37].

Chemistry and Materials Science

Factual accuracy in the chemical sciences involves not just the comprehension of text, but the spatial, structural, stoichiometric, and mathematical representation of molecules and physical properties.

ChemBench and the Paradox of Superhuman Performance

Authored by Jablonka, Ramos, and colleagues (2024), **ChemBench** provides an automated, extensive framework for evaluating the chemical knowledge, intuition, and reasoning abilities of state-of-the-art LLMs against the expertise of human chemists[8]. Comprising over 2,700 meticulously curated question-answer pairs, the benchmark goes beyond simple multiple-choice formats, requiring scientifically sound parsing that supports floating-point answers in scientific notation[8]. ChemBench employs remarkably strict accuracy thresholds, such as requiring the Mean Absolute Error (MAE) to be within 1% of the target value for numeric questions[40].

The empirical findings from ChemBench present a fascinating paradox: the best models, on average, outperformed the best human chemists in the study[8]. However, an analysis of the failure modes revealed that the models frequently struggled with fundamental, basic tasks that humans easily solve, and consistently provided highly overconfident, hallucinated predictions[8]. This dynamic—superhuman macro-performance coupled with fragile, erratic failure modes on basic logic—is a defining characteristic of current LLMs, emphasizing the need for tools that probe intuition and alignment alongside raw knowledge retrieval[39].

MACBENCH: The Multimodal Frontier

Scientific data in chemistry and materials science is inherently multimodal, heavily relying on visual spectra, microscopy, and structural diagrams. **MACBENCH** (Multimodal Materials and Chemistry Benchmark) extends the evaluation paradigm beyond text to assess Vision-Language Models (VLMs)[41]. Recognizing that earlier benchmarks focused exclusively on linguistic inputs, MACBENCH incorporates laboratory photographs, band structures, crystal systems, and atomic force microscopy (AFM) images paired with complex questions[41].

By demanding that models deduce safety hazards from photographs of laboratory apparatuses, count gold nanoislands in AFM scans, or identify crystal systems from structural diagrams, MACBENCH tests the critical intersection of visual perception and factual chemical reasoning[41]. The benchmark highlights significant gaps in the tacit knowledge embedded within current multimodal AI, proving that visual hallucinations remain a profound barrier to automated scientific assistance[41].

ChemIQ and Structural "Thinking"

Pushing beyond broad knowledge retention, **ChemIQ** focuses explicitly on molecular comprehension and advanced chemical reasoning, utilizing entirely algorithmically generated questions to ensure models cannot rely on data leakage[4]. By requiring models to interpret SMILES strings, deduce NMR elucidations, predict reactions, and perform mathematical Structure-Activity Relationship (SAR) analysis, ChemIQ demonstrates that recent reasoning-focused models (like OpenAI's o3-mini) are beginning to parse molecular graphs structurally rather than just treating them as linguistic patterns[4]. The ability of general-purpose LLMs to accurately generate IUPAC names and conduct complex reaction prediction signals a fundamental shift toward models that actually "think" about spatial and chemical logic, although perfection remains elusive outside of rudimentary tasks like ring counting[4].

Structured Knowledge and Knowledge Graphs (KGs)

Many downstream enterprise and academic applications rely on Knowledge Graphs (KGs), which store structured factual knowledge as interconnected entities and relationships (represented as triples)[42]. If an LLM is to interact with a KG for semantic search, biomedical integration, or recommendation systems, its factual accuracy must be absolute, as a single hallucinated atomic fact can corrupt the entire reasoning chain across the graph.

FactCheck and Triple Validation

Shami, Marchesin, and Silvello (2026) introduced **FactCheck**, a general-purpose benchmark designed to strictly assess LLMs in the validation of individual KG facts[42]. The pipeline systematically transforms structured triples into natural language statements and evaluates their veracity across three principal dimensions: the LLM’s internal parametric knowledge, external evidence synthesized via Retrieval-Augmented Generation (RAG), and aggregated knowledge derived from multi-model consensus strategies[42].

FactCheck distinguishes itself from existing text-based benchmarks like CRAG and FactBench by prioritizing the strict, granular logic required to validate isolated KG triples[43]. The research demonstrates that while RAG generally improves performance, high-performing QA models frequently fail at verifying atomic triples, highlighting a persistent inability of LLMs to parse structured, discrete relational logic without hallucinating contextual assumptions[42].

Evaluating Ambiguity and Source Citations

Factuality becomes incredibly complex in ambiguous scientific scenarios where multiple valid answers exist, or where competing hypotheses are debated. Patel and Anand (2024) evaluated leading models (such as GPT-4o-mini and Claude-3.5) on ambiguous QA datasets featuring source citations, including DisentQA-DupliCite and AmbigQA-Cite[3].

The evaluation revealed a stark dichotomy: while advanced models consistently predicted at least one correct answer in ambiguous contexts, they struggled significantly to handle cases requiring the acknowledgment of multiple valid answers, often defaulting to oversimplified, singular assertions[3]. More alarmingly, citation performance—the ability of the model to correctly attribute its generated facts to specific, credible sources—was exceptionally weak, with citation accuracy consistently measured at near zero across all models evaluated[3]. The introduction of *conflict-aware prompting* yielded substantial improvements, allowing models to navigate ambiguity without sacrificing factual grounding, yet the inability of baseline models to accurately cite sources remains a major obstacle to scientific trustworthiness[3].

Current Approaches and Developments in Factuality Evaluation

As the complexity of evaluating LLMs scales alongside their capabilities, the reliance on manual human evaluation has become a major bottleneck due to prohibitive costs, time constraints, and the requisite domain expertise. Consequently, recent developments have focused on creating automated, scalable frameworks capable of judging factuality with human-level precision[15].

FActScore: Fine-grained Atomic Evaluation

When evaluating long-form generated text, binary judgments (true/false) are entirely inadequate, as a single generated paragraph may contain a complex mixture of supported facts, partial truths, and outright hallucinations[15]. To resolve this evaluation crisis, Min et al. (2023) developed **FActScore** (Factual precision in Atomicity Score)[16].

FActScore systematically decomposes a model's long-form generation into a series of discrete, independent "atomic facts"[15]. Each atomic fact is then individually verified against a trusted knowledge source (such as a Wikipedia dump or a specialized scientific corpus) utilizing a strong language model acting as a judge in conjunction with a retriever[44]. The final score represents the percentage of atomic facts explicitly supported by the evidence.

This methodology provides a highly granular view of factual precision. In extensive evaluations of generated biographies, for instance, FActScore revealed that while the text appeared highly fluent and authoritative, models like ChatGPT achieved only a 58% accuracy rate on their atomic claims[15]. Furthermore, the automated FActScore estimator demonstrated an error rate of less than 2% compared to human expert annotators[15]. By automating this process, the researchers evaluated 6,500 generations from 13 different models—an endeavor that would have cost approximately $26,000 if executed by human raters—proving that highly accurate, automated factual benchmarking is achievable at scale[15].

Table 2: Analytical Metrics for Factual Evaluation
Metric Name

FActScore
Matthews Correlation Coefficient (MCC)
Unknown Rate (UnR)
Explanation Quality (EQ)

FELM and RealFactBench

The **FELM** benchmark (Benchmarking Factuality Evaluation of Large Language Models, Chen et al., 2023) provides a standardized collection of datasets specifically designed to quantify how well automated fact-checkers perform, addressing the recursive, "chicken-and-egg" problem of "evaluating the evaluators"[45].

Similarly, **RealFactBench** evaluates factuality utilizing a multi-dimensional matrix of metrics across 6,000 high-quality claims drawn from authoritative sources[49]. Beyond calculating the Macro F1-Score for factual accuracy and the Matthews Correlation Coefficient (MCC) for prediction reliability, RealFactBench introduces the concept of the **Unknown Rate (UnR)**. The UnR explicitly quantifies the frequency with which a model outputs "Unknown" when faced with insufficient evidence or conflicting information[49]. A high UnR indicates a model's capacity for caution and a resistance to hallucination, but may also signal over-conservatism, highlighting the delicate trade-off models must navigate between truthfulness and utility in real-world fact-checking[49]. Explanation Quality (EQ) is also rigorously measured using an LLM-as-a-Judge framework to assess logical coherence and evidence sufficiency, ensuring models can provide human-comprehensible justifications for their scientific claims[49].

TruthfulQA and the Phenomenon of Inverse Scaling

Lin, Hilton, and Evans (2022) introduced **TruthfulQA**, a benchmark specifically designed to measure how models mimic human falsehoods[13]. The benchmark comprises 817 questions spanning 38 categories (including health, law, and science) that target common human misconceptions, superstitions, and false beliefs[13].

The critical, counter-intuitive discovery derived from TruthfulQA is the phenomenon of **inverse scaling**. In traditional NLP tasks, performance scales positively with model size (parameter count) and training compute. However, Lin et al. observed that larger models (e.g., the 175B parameter variant of GPT-3) were systematically *less* truthful on these specific questions than their smaller counterparts[14]. Because larger models are exponentially better at capturing the statistical distribution of the pre-training data—which is heavily saturated with human errors, myths, and internet misconceptions—they are significantly more likely to generate "imitative falsehoods"[14]. While a human might achieve 94% on this benchmark by relying on verified facts, the best language model at the time achieved only 58%, smoothly synthesizing conflicting claims into highly confident, deceptive truth-statements[13]. This development proves that simply scaling up parameters and compute is insufficient for improving truthfulness; specialized alignment techniques, non-imitative training objectives, and rigorous fact-checking pipelines are strictly necessary to force models to prioritize reality over statistical likelihood[13].

Research Challenges and Limitations

Despite rapid advancements in benchmarking methodologies, several profound challenges continue to limit the efficacy of factual evaluation in scientific domains.

1. **Data Contamination and Memorization:** The most pervasive threat to benchmarking validity is the unintentional inclusion of evaluation datasets within the massive, opaque training corpora of proprietary LLMs[25]. If a model has memorized the answers to a benchmark (such as PubMedQA or GPQA) during pre-training, it simulates reasoning capabilities it does not actually possess, severely inflating perceived factual accuracy. While dynamic datasets (like SciEval) attempt to mitigate this by generating novel questions at test time, dynamically generating complex, open-ended scientific scenarios at scale remains computationally difficult and resource-intensive[25].
2. **The Limitations of Retrieval-Augmented Generation (RAG):** RAG is widely utilized to ground LLM outputs in verified external knowledge, theoretically eliminating extrinsic hallucinations by forcing the model to cite retrieved documents. However, research indicates that high-performing QA models often fail to strictly adhere to the logic of retrieved texts. Models exhibit *context-neglecting hallucinations*, where they ignore the provided context in favor of their own flawed parametric priors, or struggle to synthesize the multi-hop logical deductions required to apply retrieved chemical or physical data to a novel problem[6].
3. **Ambiguity and Multi-Truth Scenarios:** Scientific research frequently involves highly ambiguous scenarios where empirical consensus has not yet been reached, where multiple valid hypotheses exist, or where data is contradictory. Current LLMs perform exceptionally poorly in such contexts, routinely collapsing nuanced scientific debates into simplified, singular assertions[3]. Furthermore, their inability to generate accurate, traceable citations for competing theories severely limits their utility as research assistants[3].
4. **Multimodal Hallucination:** As models transition to processing visual scientific data (e.g., via MACBENCH), they inherit entirely novel vectors for hallucination. Vision-Language Models (VLMs) may misinterpret spatial relationships in chemical structures, invent features in microscopic images, or misread data visualizations. These errors require entirely new taxonomies, evaluation metrics, and benchmarking frameworks to detect, as text-based evaluation is blind to visual misinterpretation[2].

Research Gaps and Future Directions

The limitations of current benchmarking paradigms illuminate several critical trajectories for future research and development.

1. Scalable Oversight and AI-Assisted Verification: As demonstrated by the GPQA benchmark, frontier models are beginning to operate in highly specialized scientific domains where non-expert humans cannot reliably verify the output, even with unlimited internet access[29]. Future research must prioritize scalable oversight mechanisms. This involves developing frameworks where auxiliary, specialized AI models act as adversarial debaters, critiquers, or recursive reward modelers, assisting human experts in identifying deeply embedded logical flaws or subtle factual hallucinations within highly technical outputs[29].

2. Dynamic and Evolving Benchmarks: To permanently solve the data contamination issue, static benchmarks must be increasingly replaced by continuously evolving, procedurally generated evaluation suites. Drawing on the principles established by ChemIQ and the dynamic subsets of SciEval, future frameworks should leverage programmatic generation of scientific problems—such as dynamically generating novel molecular graphs, altering physical constants in physics simulations, or synthesizing artificial clinical case studies—where the ground truth is computationally verified on-the-fly rather than scraped from human literature[4].

3. Geometric and Representation-Space Detection: Traditional hallucination detection relies on analyzing the final text output of the model. A highly promising, nascent direction involves "white-box" evaluation—analyzing the internal geometric property of the LLM's representation space during generation. Techniques such as Contrastive Activation Addition (CAA) or Prompt Embedding Probes attempt to identify the literal "direction of falsehood" in the residual stream during inference, allowing for real-time hallucination detection and mitigation before the deceptive tokens are even generated[7].

4. Epistemological Grounding via Agentic Workflows: To address citation failures and multi-truth ambiguity, evaluation must shift from assessing single-pass model generation to evaluating *agentic workflows*. Benchmarks will increasingly need to measure how well an LLM can act as an autonomous agent—utilizing external tools such as Python interpreters for mathematical validation, automated literature search APIs, and database querying—to arrive at a verified, reproducible conclusion, thereby reflecting the true, iterative nature of the scientific method[8].

Conclusion

The endeavor to benchmark the factual accuracy of Large Language Models across scientific domains represents one of the most critical epistemological and technical challenges in contemporary artificial intelligence. As these models rapidly transition from conversational novelties to integral computational components of medical diagnosis, materials discovery, and quantitative academic research, the tolerance for hallucination approaches absolute zero.

The systematic review of current literature and datasets—ranging from the atomic precision of FActScore and the cognitive gradations of SciEval, to the expert-level demands of Med-PaLM

2 and the Google-proof rigor of GPQA—reveals a highly complex landscape. While models routinely demonstrate unprecedented, and often superhuman, performance on macro-level evaluations, their reliability is frequently undermined by intrinsic falsehoods, catastrophic failures on dynamic mathematical calculations, and the alarming phenomenon of inverse scaling, where enhanced statistical mimicry breeds highly persuasive deception rather than grounded truth.

Advancing the frontiers of scientific AI requires moving decisively beyond static, lexical-based evaluation. The path forward demands dynamic, algorithmically generated benchmarking to prevent memorization, the integration of rigorous multimodal evaluation frameworks to capture the visual nature of empirical science, and the rapid maturation of scalable oversight mechanisms that empower humans to verify hyper-advanced computational outputs. Ultimately, achieving reliable factual accuracy is not merely an engineering challenge of increasing parameter counts or computing power; it is an epistemological challenge requiring robust, transparent, and strictly domain-aligned architectures capable of genuine, verifiable reasoning.

Works cited

1. Towards Expert-Level Medical Question Answering with Large Language Models - arXiv, https://arxiv.org/pdf/2305.09617
2. A Survey on Hallucination in Large Language and Foundation Models - Preprints.org, https://www.preprints.org/manuscript/202504.1236
3. Factuality or Fiction? Benchmarking Modern LLMs on Ambiguous QA with Citations - arXiv, https://arxiv.org/html/2412.18051v1
4. Assessing the Chemical Intelligence of Large Language Models - arXiv, https://arxiv.org/html/2505.07735v1
5. Towards Expert-Level Medical Question Answering with Large Language Models - arXiv, https://arxiv.org/abs/2305.09617
6. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions | alphaXiv, https://www.alphaxiv.org/audio/2311.05232v1
7. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions | Request PDF - ResearchGate, https://www.researchgate.net/publication/386002726_A_Survey_on_Hallucination_in_Large_Language_Models_Principles_Taxonomy_Challenges_and_Open_Questions
8. [2404.01475] Are large language models superhuman chemists? - arXiv, https://arxiv.org/abs/2404.01475
9. A Survey of Hallucination in Large Foundation Models - Semantic Scholar, https://www.semanticscholar.org/paper/A-Survey-of-Hallucination-in-Large-Foundation-Rawte-Sheth/71bc0c97c20fffce796a355b16bd202987260029
10. A Survey on Hallucination in Large Language Models: Principles, Taxonomy,

Challenges, and Open Questions - Semantic Scholar, https://www.semanticscholar.org/paper/A-Survey-on-Hallucination-in-Large-Language-Models%3A-Huang-Yu/1e909e2a8cdacdcdff125ebcc566f37cb869a1c8
11. Large Language Models in Medicine: Applications, Challenges, and Future Directions - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC12163604/
12. TruthfulQA: Measuring How Models Mimic Human Falsehoods - ACL Anthology, https://aclanthology.org/2022.acl-long.229/
13. arXiv:2109.07958v2 [cs.CL] 8 May 2022, https://arxiv.org/pdf/2109.07958
14. FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation - ACL Anthology, https://aclanthology.org/2023.emnlp-main.741/
15. [PDF] FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation | Semantic Scholar, https://www.semanticscholar.org/paper/FActScore%3A-Fine-grained-Atomic-Evaluation-of-in-Min-Krishna/bd5deadc58ee45b5e004378ba1d54a96bc947b4a
16. Graduate-Level Google-Proof Q&A (GPQA) Benchmark - GM-RKB - Gabor Melli, http://www.gabormelli.com/RKB/Graduate-Level_Google-Proof_Q%26A_(GPQA)_Benchmark
17. PubMedQA: A Dataset for Biomedical Research Question Answering - ACL Anthology, https://aclanthology.org/D19-1259/
18. [1909.06146] PubMedQA: A Dataset for Biomedical Research Question Answering - arXiv, https://arxiv.org/abs/1909.06146
19. [2308.13149] SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research - arXiv, https://arxiv.org/abs/2308.13149
20. SciKnowEval: Evaluating Multi-level Scientific Knowledge of Large Language Models - arXiv, https://arxiv.org/html/2406.09098v1
21. A review of large language models and autonomous agents in chemistry | Chemical Science, https://pubs.rsc.org/sc/article/16/6/2514/869535/A-review-of-large-language-models-and-autonomous
22. TruthfulQA: Measuring How Models Mimic Human Falsehoods - ResearchGate, https://www.researchgate.net/publication/354650393_TruthfulQA_Measuring_How_Models_Mimic_Human_Falsehoods
23. SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research, https://www.alphaxiv.org/abs/2308.13149v2
24. [AAAI 2024] SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research · GitHub, https://github.com/OpenDFM/SciEval
25. SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research, https://www.researchgate.net/publication/373437733_SciEval_A_Multi-Level_Large_Language_Model_Evaluation_Benchmark_for_Scientific_Research
26. SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research, https://arxiv.org/html/2308.13149v2
27. GPQA: A Graduate-Level Google-Proof Q&A Benchmark | alphaXiv, https://www.alphaxiv.org/abs/2311.12022

28. GPQA: A Graduate-Level Google-Proof Q&A Benchmark | OpenReview, https://openreview.net/forum?id=Ti67584b98
29. SCIBENCH - OpenDataLab, https://opendatalab.com/OpenDataLab/SCIBENCH
30. (PDF) Toward expert-level medical question answering with large language models, https://www.researchgate.net/publication/387826586_Toward_expert-level_medical_question_answering_with_large_language_models
31. (PDF) Towards Expert-Level Medical Question Answering with Large Language Models, https://www.researchgate.net/publication/370814210_Towards_Expert-Level_Medical_Question_Answering_with_Large_Language_Models
32. PubMedQA: A Dataset for Biomedical Research Question Answering - ResearchGate, https://www.researchgate.net/publication/335833303_PubMedQA_A_Dataset_for_Biomedical_Research_Question_Answering
33. PubMedQA Homepage, https://pubmedqa.github.io/
34. Benchmarking Large Language Models from Open and Closed Source Models to Apply Data Annotation for Free-Text Criteria in Healthcare - MDPI, https://www.mdpi.com/1999-5903/17/4/138
35. Med-HALT: Medical Domain Hallucination Test for Large Language Models, https://www.semanticscholar.org/paper/Med-HALT%3A-Medical-Domain-Hallucination-Test-for-Umapathi-Pal/3b0792f6d7f6aa6aadd316e73943116afef2979b
36. Med-HALT: Medical Domain Hallucination Test for Large Language Models - ResearchGate, https://www.researchgate.net/publication/372767219_Med-HALT_Medical_Domain_Hallucination_Test_for_Large_Language_Models
37. (PDF) A framework for evaluating the chemical knowledge and reasoning abilities of large language models against the expertise of chemists - ResearchGate, https://www.researchgate.net/publication/391910567_A_framework_for_evaluating_the_chemical_knowledge_and_reasoning_abilities_of_large_language_models_against_the_expertise_of_chemists
38. chembench/eval-card.md at dev - GitHub, https://github.com/lamalab-org/chembench/blob/dev/eval-card.md
39. MACBENCH: A multimodal chemistry and materials science benchmark - OpenReview, https://openreview.net/pdf/ae3341fda55515c88b8e83fc192f62b4d129e913.pdf
40. Benchmarking Large Language Models for Knowledge Graph Validation - arXiv, https://arxiv.org/html/2602.10748v1
41. Benchmarking Large Language Models for Knowledge Graph Validation - OpenProceedings.org, https://openproceedings.org/2026/conf/edbt/paper-187.pdf
42. FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation - arXiv, https://arxiv.org/html/2305.14251v2
43. [2310.00741] FELM: Benchmarking Factuality Evaluation of Large Language Models - arXiv, https://arxiv.org/abs/2310.00741

44. FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation | Request PDF - ResearchGate, https://www.researchgate.net/publication/370981225_FActScore_Fine-grained_Atomic_Evaluation_of_Factual_Precision_in_Long_Form_Text_Generation
45. factscore - PyPI, https://pypi.org/project/factscore/
46. [2305.14251] FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation - arXiv, https://arxiv.org/abs/2305.14251
47. RealFactBench: A Benchmark for Evaluating Large Language Models in Real-World Fact-Checking - arXiv, https://arxiv.org/html/2506.12538v1
48. NeurIPS 2023 Track Datasets and Benchmarks Submissions - OpenReview, https://openreview.net/submissions?page=2&venue=NeurIPS.cc%2F2023%2FTrack%2FDatasets_and_Benchmarks
49. TruthfulQA: Measuring How Models Mimic Human Falsehoods - ResearchGate, https://www.researchgate.net/publication/361063493_TruthfulQA_Measuring_How_Models_Mimic_Human_Falsehoods
50. A Survey on Hallucination in Large Vision-Language Models - arXiv, https://arxiv.org/html/2402.00253v2